SKILLNAVORA · AI PRODUCT MANAGEMENT & GENAI LAUNCHPAD

AI Product Strategy

Decide where AI creates real value: AI-first vs AI-added products, moats and data flywheels, build vs buy vs partner, model choice, opportunity assessment and AI roadmaps.

Edition 2026 · skillnavora.in

How to use this guide

Each class has goals, concepts with Indian examples, key takeaways and a practice exercise, from basic to advanced. Finish with the final project to build something you can show in interviews.

	Class	Level
Class 1	**Where AI creates value**	Basic
Class 2	**Moats and data flywheels**	Basic
Class 3	**Build vs buy vs partner, and choosing models**	Intermediate
Class 4	**Assessing and prioritising AI opportunities**	Intermediate
Class 5	**Market and competitor analysis for AI products**	Intermediate
Class 6	**The business case: cost modelling, ROI and pilots**	Intermediate
Class 7	**AI roadmaps and portfolio of bets**	Advanced
Class 8	**The AI operating model: teams, skills and governance**	Advanced
Class 9	**AI strategy by industry**	Advanced
Class 10	**Writing and communicating the AI strategy**	Advanced
Class 11	**Case studies: three AI strategy decisions**	Advanced
Class 12	**Interview practice: AI strategy scenario questions**	Advanced
Class 13	**Case studies and final project**	Advanced
	Toolkit: templates, cheat sheet, glossary, 16 interview questions, further reading	

BASIC · CLASS 1

Where AI creates value

Separate real AI opportunities from hype.



AI opportunity matrix

AI-first vs AI-added

AI-added	AI-first
AI improves an existing workflow (summaries in a CRM, smart replies)	The product only exists because of AI (AI tutor, coding agent, voice receptionist)
Lower risk, faster	Bigger upside, more uncertainty

Good AI opportunities usually have

- A frequent, painful job with lots of language, images, or decisions.
- Tolerance for some errors, or a cheap human check.
- Data or context you can access (documents, tickets, history).
- A clear metric that improves (time saved, conversion, resolution).

Worked example: where would AI help a grocery delivery app?

Idea	User value	Feasibility today	Error tolerance	Verdict
Smart substitutions when an item is out of stock	High: fewer cancelled orders	High (rules + embeddings)	Medium: wrong swap can be refunded	Quick win
Chat assistant that plans a week's meals and fills the cart	High for busy users	Medium (needs catalogue grounding)	Medium: user reviews cart	Big bet, pilot first
AI that auto-approves refunds	Medium	High	Low: money and fraud risk	Human in the loop, strict limits
AI-written product descriptions for 50,000 SKUs	Medium	High	High if reviewed	Quick win with sampling checks
Autonomous pricing of every item	Unclear	Medium	Very low: legal and trust risk	Avoid for now

Rule of thumb: start where the user already feels the pain, the data already exists, a wrong answer is cheap to fix, and success is measurable within a month.

Key takeaways

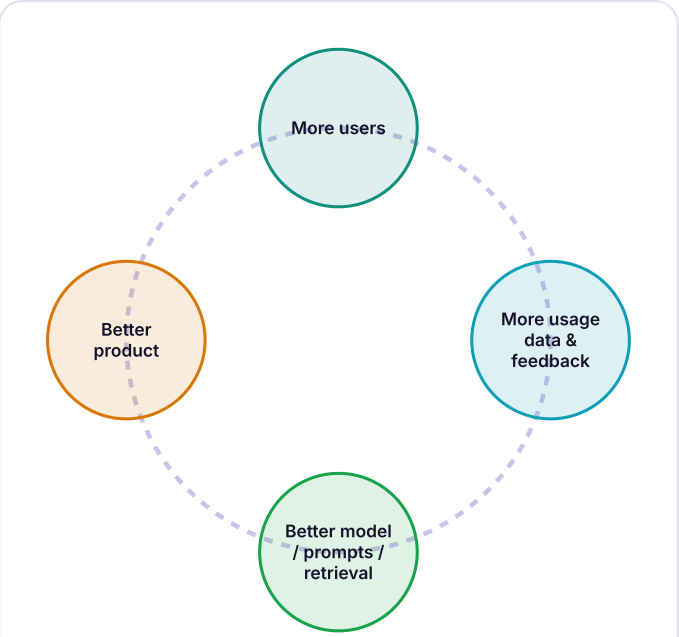
- Start from the user job, not the model.
- Plot value vs feasibility.
- Error tolerance decides how much autonomy AI gets.

Practice exercise: List 10 AI ideas for a product you know and place them on the opportunity matrix.

BASIC · CLASS 2

Moats and data flywheels

Understand what makes an AI product hard to copy.



The data flywheel

Where advantage comes from (when everyone can call the same models)

- **Proprietary data and context**: your customers' history, domain documents, labelled outcomes.
- **Workflow ownership**: being where the work happens (the system of record).
- **Distribution**: existing users, brand, partnerships.
- **Trust and quality**: evals, safety, reliability in a domain (health, finance, legal).
- **Speed of learning**: feedback loops that improve weekly.

> Thin "wrapper" products that only add a UI on a public model are easy to copy; add data, workflow, and trust.

Five moats, ranked by how hard they are to copy

Moat	Example	How it erodes	How to build it
Proprietary data and feedback loops	Years of resolved support tickets and outcomes	Competitors buy or generate similar data	Capture outcome labels (was the answer accepted?), not only inputs
Workflow ownership / system of record	The CRM where reps log every deal	Platform bundles the feature	Be where work starts and finishes; integrate deeply
Distribution and trust	A bank's customer base	New channels, new brands	Use the base to learn faster than outsiders
Domain expertise baked into the product	Medical coding rules, tax logic	Experts are hired away	Encode expertise as evals, templates, guardrails
Cost and speed advantage	Smaller, tuned models for your task	Model prices fall for everyone	Treat it as temporary; reinvest savings

Wrapper risk: if your only value is a nicer chat window on top of a public model, a model provider or a competitor can copy it in weeks. Ask: *what do we know, own or control that a new entrant cannot get in six months?*

Key takeaways

- Models are commodities; data, workflow and trust are moats.
- Design feedback loops from day one.
- Measure the flywheel: does quality improve with usage?

Practice exercise: Draw the data flywheel for an AI feature you'd build and name the metric for each step.

INTERMEDIATE · CLASS 3

Build vs buy vs partner, and choosing models

Make sourcing and model decisions with clear trade-offs.

Option	When	Risks
Buy (SaaS AI tool)	Non-core need, fast time to value	Vendor lock-in, data sharing, less differentiation
Build on model APIs	Core experience, need control over prompts, data, UX	Team skills, evals, cost management
Partner	Specialised data or distribution you lack	Dependency, revenue share
Fine-tune / self-host open models	Strict data residency, high volume, specialised tasks	Infrastructure, maintenance, quality gap

Model choice criteria

- Quality on **your** eval set (not public leaderboards alone).
- Cost per successful task; latency; context length.
- Data policies, residency (India data needs), enterprise terms.
- Multilingual quality (Hindi and regional languages), safety behaviour.
- Keep a provider-agnostic layer so you can switch.

A model-choice scorecard

Criterion	Question	Weight (example)
Quality on our task	Score on our own eval set, not public leaderboards	30%
Cost per task	Tokens x price, retries, tools	20%
Latency	Time to first token and to full answer	15%
Privacy and residency	Where data goes, retention, India hosting, DPAs	15%
Reliability and limits	Uptime, rate limits, support, deprecation policy	10%
Lock-in	How hard to switch (prompts, tool formats, fine-tunes)	10%

- **Multi-model routing:** a small cheap model for easy requests, a large model for hard ones, with an eval-driven fallback.
- **Open-weight models:** run in your cloud for privacy or cost at scale; you take on hosting, security and upgrades.
- **Contract checks:** training on your data (should be off), data retention, IP indemnity, SLAs, exit terms.

Key takeaways

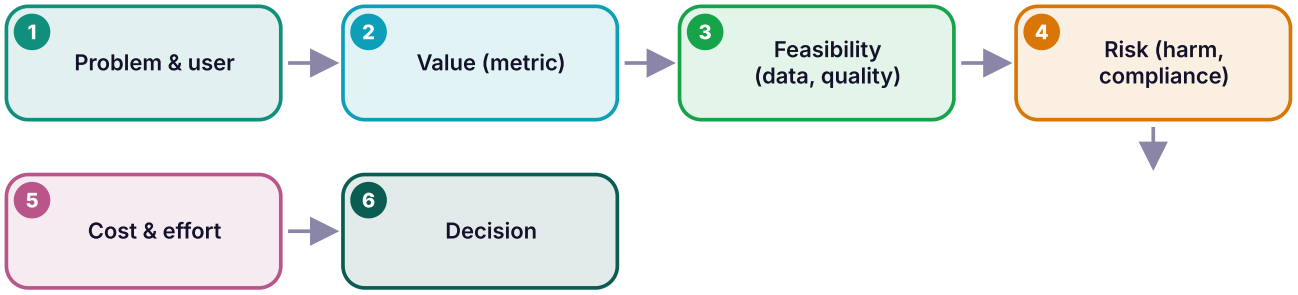
- Buy non-core, build core, partner for missing assets.
- Choose models with your own evals and cost per task.
- Avoid lock-in with an abstraction layer.

Practice exercise: Write a build-vs-buy memo for an AI support assistant at a 200-person company.

INTERMEDIATE · CLASS 4

Assessing and prioritising AI opportunities

Score AI ideas on value, feasibility and risk.



AI opportunity assessment

Question	Evidence
How often and how painful is the job?	Usage data, interviews
Can AI do it well enough today?	Prototype + 30-case eval
What happens when it's wrong?	Failure modes, human check cost
Cost per task vs value per task	Token estimates, pricing
Data access and permissions	Legal and security review

Run a **2-week feasibility spike**: prototype, small eval, cost estimate, user reaction, before committing roadmap space.

Scoring template (use 1 to 5)

Opportunity	User impact	Business impact	Feasibility	Risk (reverse)	Data readiness	Total
Smart substitutions	5	4	4	4	4	21
Meal-planning assistant	4	4	3	3	3	17
Auto-refunds	3	3	4	1	4	15

Then ask three kill questions: **Is there a cheaper non-AI solution? What does a wrong answer cost? Can we measure value in 4 weeks?**

Key takeaways

- Prototype + eval before roadmap commitment.
- Compare cost per task with value per task.
- Price in the cost of errors.

Practice exercise: Assess 3 AI ideas with the table above and recommend one.

INTERMEDIATE · CLASS 5

Market and competitor analysis for AI products

Understand the landscape before choosing where to play.



Where AI features sit in the market

AI markets move fast and features become table stakes within months. Strategy starts with an honest map of who does what.

Player type	Strengths	Weaknesses	Your response
Big platforms (Google, Microsoft, Salesforce, Amazon)	Distribution, data, bundling	Slow, generic	Be deeper in a niche workflow; integrate rather than fight
Model labs (OpenAI, Anthropic, Google DeepMind, Meta)	Best models, fast release cadence	Horizontal, not workflow-specific	Build on them; avoid pure wrappers
AI-native startups	Speed, focus, new UX	Little data and trust, funding risk	Learn from their UX; compete on trust and integration
Incumbents in your industry	Customers and regulation know-how	Legacy systems	Move faster; partner where useful
Open-source communities	Low cost, flexibility	Support and polish	Use for components; contribute selectively

Competitor teardown template

- **Job:** what user job do they solve, for whom?
- **Experience:** screenshots of the best 5 moments and the worst 3.
- **Tech clues:** model used, latency, grounding, price per seat or per use.
- **Proof:** reviews, case studies, funding, hiring (what roles are they posting?).
- **Gaps:** what users complain about; which segments they ignore.

> Track the **release cadence** of the three largest models and platforms in your space. When they ship a feature you rely on, you must already know what you will do next.

Key takeaways

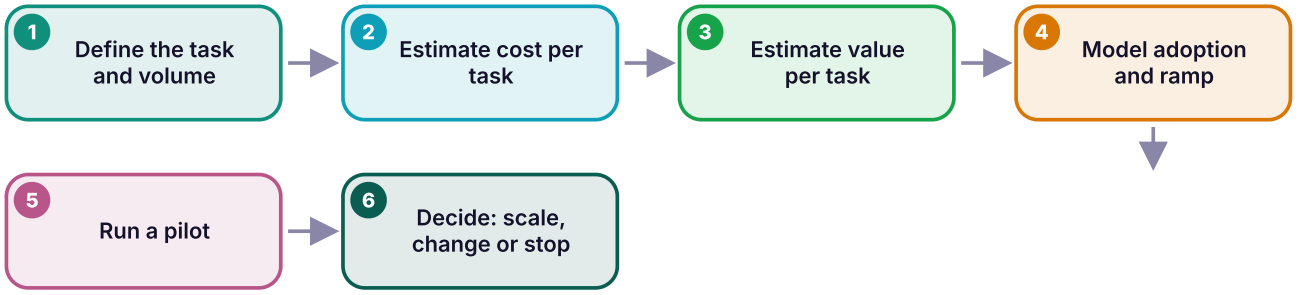
- Map players by type and response, not only by name.
- Table stakes cannot be a strategy.
- Tear down competitors on job, experience, tech and gaps.

Practice exercise: Teardown two AI competitors with the template and write a 5-line positioning statement.

INTERMEDIATE · CLASS 6

The business case: cost modelling, ROI and pilots

Show leaders the numbers for an AI feature.



From idea to investment decision

Cost per task (example: a support-reply assistant)

Item	Assumption	Cost per ticket
Model usage	3,000 input + 500 output tokens, one retry in 10%	Rs 0.60
Retrieval and tools	Vector search, helpdesk API calls	Rs 0.10
Evaluation and monitoring	Allocated across volume	Rs 0.15
Human review (20% of tickets)	2 minutes at Rs 6/min	Rs 2.40
Total		about Rs 3.25

Value per ticket: a human reply costs about Rs 25 and takes 8 minutes. If AI drafts handle 60% of tickets with edits of 2 minutes, saving is about Rs 17 per eligible ticket. At 100,000 tickets a month: 60,000 x Rs 17 = about **Rs 10 lakh saved** against about Rs 3 lakh of run cost, plus build cost. These numbers are illustrative; use your own.

Total cost of ownership (TCO)

- **Build:** engineering, design, data work, evals, legal review.
- **Run:** inference, vector database, monitoring, human review, support.
- **Change:** prompt and model updates, re-evaluation, retraining, vendor price changes.
- **Risk reserve:** incidents, compliance, brand damage.

Pilot design

- One workflow, one team, 4 to 8 weeks, a control group, pre-agreed success metrics (time saved, quality score, CSAT, cost per task).
- Stop rule: define in advance what result means "stop" or "change", so politics does not decide.

Key takeaways

- Model cost per task, not per month.
- Include human review and monitoring in TCO.
- Decide with pre-agreed success and stop rules.

Practice exercise: Build a cost and value model for an AI feature in your product with 3 scenarios (low, base, high adoption).

ADVANCED · CLASS 7

AI roadmaps and portfolio of bets

Plan AI work under uncertainty.



Three horizons for AI investment

- Use outcome-based roadmaps with confidence levels; AI capability changes fast.
- Budget for evals, data work and safety, not only features.
- Revisit quarterly: new models can make "hard" ideas easy (and cheaper).
- Kill criteria: what eval or adoption level stops a bet.

Portfolio of bets (70-20-10)

Bucket	Share of effort	What goes here	Success looks like
Core improvements	70%	AI that improves existing flows (search, support, ops)	Clear metric lift, low risk
Adjacent bets	20%	New AI features for current users	Adoption and retention uplift
Transformational	10%	AI-first products or workflows	Learning, option value, a few strong signals

Revisit quarterly. Kill bets that miss learning goals; double down where usage and retention show up.

Key takeaways

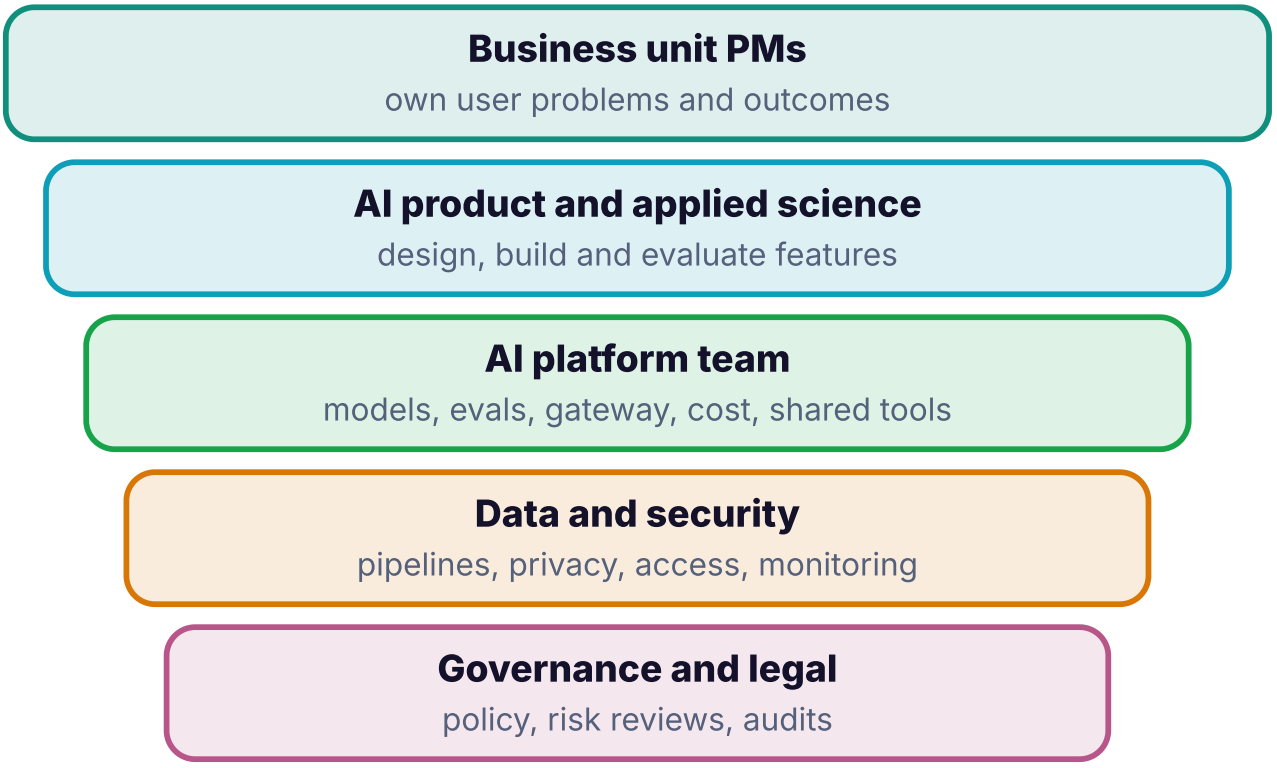
- Balance quick wins with bigger bets.
- Include evals, data and safety work in the plan.
- Set kill criteria for each bet.

Practice exercise: Create a 3-horizon AI roadmap for an Indian fintech app.

ADVANCED · CLASS 8

The AI operating model: teams, skills and governance

Organise people so AI products ship safely and learn fast.



A typical AI operating model

Role	Does
AI product manager	Defines problem, evals and trade-offs; owns outcomes and risk acceptance
ML / applied scientist	Model choice, fine-tuning, experiments, error analysis
AI / software engineer	Pipelines, retrieval, tool integrations, latency and cost
Data engineer	Quality data, labelling flows, feature stores
Designer / researcher	AI UX, trust, feedback loops (see the AI UX course)
Domain expert	Reviews outputs, writes golden examples
Trust, security and legal	Risk reviews, policy, incident response

Centralised, embedded or hybrid?

Model	Good for	Risk
Central AI team	Early stage, scarce skills	Bottleneck, far from users
Embedded in product teams	Speed and domain fit	Duplicated effort, uneven quality
Hybrid (platform + embedded)	Scale	Needs clear interfaces and ownership

Rule: share the platform (models, evals, gateway, guardrails), embed the product work, and keep one accountable owner for each AI feature in production.

Key takeaways

- Hybrid: shared platform plus embedded product teams.
- Each AI feature has one accountable owner.
- Governance is part of the operating model, not an afterthought.

Practice exercise: Draw the team and ownership model for an AI assistant in your organisation.

ADVANCED · CLASS 9

AI strategy by industry

See where AI creates value, and where to be careful, in different sectors.

Industry	High-value use cases	Be careful with	Key metric
E-commerce and retail	Search and discovery, substitutions, support, catalogue enrichment, demand forecasting	Price manipulation claims, fake reviews, personalisation creepiness	Conversion, AOV, return rate, support cost
Fintech and banking	Fraud detection, collections, KYC automation, advisory assistants, support	Credit decisions (fairness, explainability), advice rules, data localisation	Fraud loss, approval rate, time to onboard
Healthcare	Triage assist, documentation, imaging support, patient communication	Clinical safety, consent, hallucination, regulation	Clinician time saved, error rate
Education	Tutors, assessment help, content creation, personalised practice	Accuracy, cheating, child data	Learning outcomes, completion
B2B SaaS	Copilots in workflow, summarisation, automation	Data leakage across customers, pricing of compute	Activation, retention, expansion revenue
Logistics and mobility	ETA prediction, routing, demand matching, document processing	Real-time reliability, safety	On-time rate, cost per delivery
Media and marketing	Content generation, localisation, personalisation	Copyright, brand safety, deepfakes	Output per person, engagement

For each industry, repeat the same questions: what job, what data, what error cost, what regulation, what moat?

Key takeaways

- Industry context changes the risk, not the method.
- High-stakes domains need human review and strong evals.
- Always tie AI to an industry metric.

Practice exercise: Pick one industry and write the top three AI opportunities with metric and main risk for each.

ADVANCED · CLASS 10

Writing and communicating the AI strategy

Turn analysis into a document and story that leaders can approve.

One-page AI strategy template

Section	Content
Vision	What the product looks like in 2 to 3 years with AI
Where to play	Users, jobs, and the 3 to 5 use cases chosen (and rejected)
How to win	Moats: data, workflow, trust, distribution, cost
Capabilities needed	Data, models, platform, people, partners
Roadmap	Now (0 to 3 months), next (3 to 9), later (9+), with bets
Metrics and OKRs	North star, adoption, quality, cost, risk
Risks and guardrails	Top risks, owners and mitigations
Investment ask	Budget, headcount, timeline, decision date

Communicating under uncertainty

- Say what is known, what is assumed and what a pilot will tell you.
- Use ranges and scenarios, not single numbers.
- Show how fast you can change course (model swaps, kill criteria).
- Prepare answers for the usual executive questions: "Why now?", "What if a model provider copies this?", "What if it is wrong?", "What does it cost at 10x?", "What do we stop doing?"

> The best AI strategy documents are short, make clear choices, and say what the company will not do.

Key takeaways

- One page, clear choices, explicit non-goals.
- Show assumptions, ranges and kill criteria.
- Prepare for the five standard executive questions.

Practice exercise: Write your one-page AI strategy for a product you know and rehearse the five executive questions.

ADVANCED · CLASS 11

Case studies: three AI strategy decisions

Practise the reasoning on realistic situations (illustrative numbers).

Case 1: A fintech wants an AI advisor

Situation: 2 million users, 4% invest. Leadership wants a chatbot advisor in 90 days.

Step	Reasoning
Job	Help first-time investors understand and start, not recommend specific securities (regulated)
Data	Educational content, product catalogue, anonymised funnel data
Risk	Mis-selling; suitability rules; hallucinated returns
Decision	Launch a **guided education and goal-planning assistant** with citations and human handoff; no personalised stock picks
Metric	Start-investing conversion, drop-off at KYC, complaints, assistant accuracy on an eval set
Moat	Own funnel data and trust; integrate in the onboarding flow

Case 2: A marketplace faces model-provider dependency

Situation: every AI feature calls one model API; price rises 30%.

- Create a model gateway; route easy tasks to a cheaper model; keep evals to switch providers; negotiate volume pricing; cache.
- Result target: cost per task down 40% with quality within 2 points of baseline.

Case 3: A B2B SaaS leader says 'copy the competitor's copilot'

- Teardown showed the competitor's copilot is a thin chat over documentation.
- Choose instead an **in-workflow** assistant that drafts the three most repetitive tasks, using customer data; measure time saved per task.
- Differentiate on workflow and data, not on a chat window.

Key takeaways

- Start from the job and the risk.
- Use a gateway and evals to reduce provider risk.
- Differentiate on workflow and data.

Practice exercise: Write your own one-page decision memo for Case 1 with a rejected option.

ADVANCED · CLASS 12

Interview practice: AI strategy scenario questions

Answer strategy questions with structure and numbers.

Q. Should our company build or buy an AI support assistant?
Model answer: I would first set the goal (resolution rate and cost per ticket), then check whether support is core and differentiating. Buy for speed if it is not core, ideally a vendor that grounds on our knowledge base with exportable data. Build if we need deep integration and control. I would run a 6-week pilot on one queue with a control group, compare cost per resolved ticket, CSAT and escalation rate, and decide with pre-agreed thresholds.
Why it works: It sets the goal, applies a build/buy test, and proposes a measurable pilot with stop rules.

Q. A competitor launched an AI feature. Do we respond?
Model answer: I would check whether customers actually care and whether it is table stakes. I would test the feature with our customers, estimate cost and moat, and decide among match, leapfrog or ignore. Matching quickly with a good-enough version is right if churn risk is real; otherwise I would invest in our own differentiator.
Why it works: It avoids reactive copying and ties the choice to customer evidence.

Q. How do you pick the first AI use case for a company?
Model answer: I list jobs with frequent pain, available data and low error cost, score them on impact, feasibility, risk and data readiness, and choose one that can show value in 4 weeks with a clear metric. A quick win builds trust and funds bigger bets.
Why it works: It shows a scoring approach and sequencing logic.

Q. What would make you stop an AI project?
Model answer: Pre-agreed kill criteria: quality below threshold after two iterations, cost per task above the value created, unresolved safety or compliance risk, or low adoption despite fixes.
Why it works: Strong PMs define exit criteria before starting.

Key takeaways

- Answer with goal, options, evidence, decision and metric.
- State stop criteria.
- Link to customers and economics.

Practice exercise: Record yourself answering all four questions in under 2 minutes each.

ADVANCED · CLASS 13

Case studies and final project

Apply the strategy tools end to end.

Mini cases

- **Customer support (B2B SaaS)**: AI-added copilot first (draft replies), then AI-first self-service; moat = resolved-ticket data + workflow integration.
- **Edtech tutor**: AI-first; moat = curriculum-aligned content, learner progress data, trust with parents.
- **Fintech**: AI for fraud and support; regulatory trust as moat; heavy human-in-the-loop.

Final project

Write a 2-page AI strategy for a product: opportunity matrix, flywheel, build/buy/partner decision, model-choice criteria, 3-horizon roadmap with kill criteria, and the evals that prove progress.

Key takeaways

- Strategy = where to play + how to win + what to build.
- Tie every bet to evals and metrics.
- Revisit as models change.

Practice exercise: Complete the 2-page AI strategy and present it in 7 minutes.

TOOLKIT

Templates

AI strategy canvas

Field	Fill in
Target users & jobs	
AI opportunities (value × feasibility)	
Moat & flywheel	
Build / buy / partner	
Model criteria	
Horizon 1 / 2 / 3 bets	
Evals & metrics	
Kill criteria	

TOOLKIT

Cheat sheet

Term / tool	Meaning / use
AI-added	AI improves an existing workflow
AI-first	Product only possible because of AI
Data flywheel	Usage → data → better product → more usage
Moat	Durable advantage competitors can't copy quickly
Feasibility spike	Short prototype + eval to test viability
Kill criteria	Pre-agreed signal to stop a bet
Value test	Frequent pain, data exists, error cost low, measurable in a month
Moat test	What can't a new entrant get in 6 months?
Portfolio	70% core, 20% adjacent, 10% transformational
Stop rule	Agree success and stop thresholds before the pilot

Glossary

Wrapper	Thin product on top of a public model
System of record	Where the main business data lives
TCO	Total cost of ownership: build, run, change and risk costs
Wrapper product	A thin interface on a public model with little defensibility
Table stakes	Features customers expect from everyone
Moat	A lasting advantage that is hard to copy
Routing	Sending each request to the cheapest model that can do it well
Open-weight model	A model whose weights you can download and run
Pilot	A time-boxed test with defined success and stop criteria
Kill criteria	Pre-agreed results that end a bet
Eval set	Examples used to score model quality on your task
Unit economics	Revenue and cost per unit of use
Adoption ramp	How usage grows after launch
AI-native	A product designed around AI from the start

INTERVIEW PREP

16 interview questions with answers

Q1. How do you decide where to use AI?
Start with frequent, painful jobs; check value, feasibility (prototype + eval), risk of errors and cost per task; plot on a value-feasibility matrix.

Q2. What is a moat for an AI product?
Proprietary data and context, workflow ownership, distribution, trust/quality in a domain, and fast learning loops; not the model itself.

Q3. Build vs buy for AI?
Buy non-core needs, build the core experience on model APIs, partner for missing data or distribution; consider lock-in, data policies and cost.

Q4. How do you choose a model?
Your own evals, cost per successful task, latency, context, data policies and residency, multilingual quality, with an abstraction layer to switch.

Q5. What is a data flywheel?
A loop where usage creates data and feedback that improves quality, attracting more usage.

Q6. How do you roadmap AI under uncertainty?
Outcome roadmaps with confidence levels, three horizons, quarterly review, and kill criteria tied to evals and adoption.

Q7. What is the difference between an AI strategy and an AI feature list?
A strategy chooses where to play and how to win (users, moats, capabilities, risks, investment). A feature list is a set of tasks. Without strategy, features get copied and quickly become table stakes.

Q8. How do you estimate the ROI of an AI feature?
Model cost per task (model, retrieval, monitoring, review) and value per task (time or money saved, conversion lift), multiply by volume and adoption, add build and change costs, and validate with a pilot.

Q9. When should you not use AI?
When a rule, search filter or simple automation solves the problem, when errors are very costly and cannot be reviewed, when data is unavailable, or when value cannot be measured.

Q10. How do you choose between models?
Score candidates on your own eval set for quality, plus cost, latency, privacy, reliability and lock-in; consider routing small models for easy tasks.

Q11. What is a data flywheel and when does it fail?
More usage produces data that improves the product, attracting more usage. It fails when feedback is not captured, data is low quality, or competitors can get equivalent data.

Q12. How do you organise an AI team?
A hybrid model: a platform team for shared models, evals and guardrails, with AI product teams embedded in business units and clear accountability for each feature.

Q13. What is the biggest strategic risk of building on third-party models?
Dependency and commoditisation: price changes, deprecations, policy changes and competitors

using the same model. Mitigate with abstraction, evals, multi-model options and proprietary data or workflow advantages.

Q14. How would you decide between build, buy and partner?
Buy for non-core needs, build where it is core and differentiating, partner for data or distribution you lack. Weigh time to value, control, cost, risk and lock-in.

Q15. How do you balance quick wins and big bets?
Use a portfolio such as 70-20-10, with explicit learning goals and kill criteria for bets.

Q16. What would you tell a CEO who says 'put AI everywhere'?
Pick three high-value use cases tied to company metrics, prove value in a pilot, build shared evals and governance, then scale.

LEARN MORE

Further reading and tools

- a16z: AI articles: https://a16z.com/ai/
- Sequoia: AI perspectives: https://www.sequoiacap.com/our-companies/
- Lenny's Newsletter: AI product: https://www.lennysnewsletter.com/
- Reforge: AI strategy: https://www.reforge.com/
- McKinsey: the state of AI (annual survey): https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai
- Stanford HAI AI Index: https://hai.stanford.edu/ai-index
- a16z: AI articles: https://a16z.com/ai/

© SkillNavora · skillnavora.in · Part of the AI Product Management and GenAI Launchpad programs. Practise with the free AI Tutor, AI PM Toolkit and Practice Lab on skillnavora.in.