SKILLNAVORA · AI PRODUCT MANAGEMENT & GENAI LAUNCHPAD

Designing AI Experiences (AI UX)

Design AI features people trust and use: interaction patterns (copilot, chat, agent, inline), setting expectations, explainability and citations, feedback and control, failure and handoff design, and testing AI UX.

Edition 2026 · skillnavora.in

How to use this guide

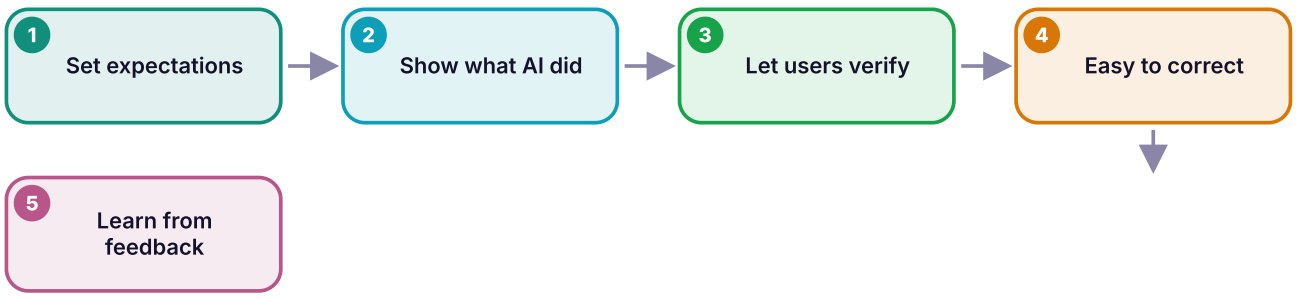
Each class has goals, concepts with Indian examples, key takeaways and a practice exercise, from basic to advanced. Finish with the final project to build something you can show in interviews.

	Class	Level
Class 1	**Principles of AI UX**	Basic
Class 2	**Interaction patterns**	Basic
Class 3	**Trust, explainability and citations**	Intermediate
Class 4	**Feedback, control and personalisation**	Intermediate
Class 5	**Conversational UI design in detail**	Intermediate
Class 6	**Designing for uncertainty, hallucination and verification**	Intermediate
Class 7	**AI UX in e-commerce and fintech: patterns and cases**	Intermediate
Class 8	**Failure, safety and handoff design**	Advanced
Class 9	**Agentic UX: autonomy, approvals and trust**	Advanced
Class 10	**Research, evaluation and metrics for AI UX**	Advanced
Class 11	**Inclusive, multilingual and voice-first AI design**	Advanced
Class 12	**Case studies: three AI experiences redesigned**	Advanced
Class 13	**Interview practice: AI UX scenario questions**	Advanced
Class 14	**Testing AI UX and final project**	Advanced
	Toolkit: templates, cheat sheet, glossary, 15 interview questions, further reading	

BASIC · CLASS 1

Principles of AI UX

Design for uncertainty and the right level of trust.



AI UX loop

- **Set expectations**: say what the AI can and can't do; show examples.
- **Calibrated trust**: not too much (blind acceptance), not too little (ignored).
- **User in control**: undo, edit, regenerate, opt out.
- **Transparency**: label AI content; show sources.
- **Graceful failure**: say "I'm not sure" and offer a human or manual path.

Ten principles for AI interfaces

Principle	In practice	Example
Set expectations	Say what the AI can and cannot do	Welcome message: 'I can help with orders, returns and delivery. I cannot change prices.'
Make the value obvious	Show a concrete first win in seconds	Suggested prompt chips tied to real tasks
Keep users in control	Preview, edit, undo, stop	Draft appears in an editable box, not sent automatically
Show your work	Sources, steps, reasons	'Based on your last 3 orders and this offer page'
Fail gracefully	Admit limits and offer a next step	'I am not sure. Connect me to an agent'
Reduce effort	Smart defaults, suggestions, shortcuts	One-tap 'shorter / friendlier' buttons
Learn from feedback	Make correction easy and visible	Thumbs with reason; 'Remember this' option
Be consistent	Same behaviour and tone	Defined voice and refusal pattern
Respect privacy	Tell what data is used	'Using your order history. Change'
Calibrate trust	Neither over- nor under-trust	Confidence cues; verify links for high-stakes answers

Key takeaways

- Design for being wrong sometimes.
- Aim for calibrated trust.
- Always provide control and recovery.

Practice exercise: Audit an AI feature you use against the five principles.

BASIC · CLASS 2

Interaction patterns

Pick the right pattern for the job.

Pattern	Example	Best for
Inline suggestion / autocomplete	Gmail Smart Compose, code completion	Small, frequent help
Copilot side panel	Docs or CRM assistant	Help while staying in the workflow
Chat	Support bot, tutor	Open questions, exploration
Generate-then-edit	Draft email, PRD, image	Creation with human polish
Agent with approvals	Booking, research agent	Multi-step tasks
Invisible AI	Ranking, fraud scores, smart defaults	Decisions users don't need to see

Chat is not always the answer; often a button ("Summarise", "Fix grammar") inside the workflow converts better.

Pattern library

Pattern	Best for	Watch out
Chat	Open-ended questions, support, exploration	Blank-page problem; long answers
Inline assist (in the editor or form)	Drafting, rewriting, autofill	Interrupting flow; unclear what changed
Command palette / slash commands	Power users, quick actions	Discoverability
Suggestions and chips	Guiding new users	Irrelevant suggestions erode trust
Side-by-side compare	Choosing between options	Complexity
Voice	Hands-free, low literacy, in-car	Noise, accents, privacy
Summaries and insights	Long content, dashboards	Hidden errors; need links to source
Agents with approvals	Multi-step tasks	Over-autonomy; see the agentic class

Key takeaways

- Match pattern to task size and risk.
- Prefer in-workflow buttons over generic chat.
- Agents need approvals and progress views.

Practice exercise: Redesign one chat feature as an inline pattern and compare.

INTERMEDIATE · CLASS 3

Trust, explainability and citations

Help users judge AI output quickly.

- Cite sources with clickable references; highlight the supporting text.
- Show confidence carefully (e.g. "Needs review" labels) rather than fake precision.
- Explain decisions in plain words ("Flagged because the amount is 5× your usual").
- Show what data was used and allow users to exclude data.
- Label AI-generated content, especially images and voice.

Citations done well

- Cite **specific passages**, not just a list of links.
- Make them **checkable in one click** (open the document at the quoted line).
- Show **when sources conflict** or are old ("Policy updated 3 months ago").
- For high-stakes answers (money, health, legal), prompt users to verify and show the official source.

Key takeaways

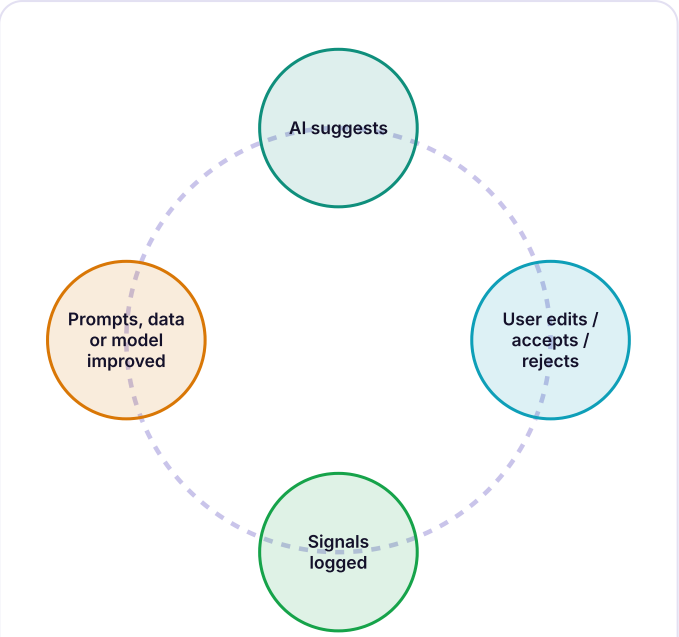
- Citations and highlights speed verification.
- Explain decisions in user language.
- Label AI content clearly.

Practice exercise: Design the citation UI for an AI answer from 3 policy documents.

INTERMEDIATE · CLASS 4

Feedback, control and personalisation

Turn user actions into learning signals.



Feedback loop

- Explicit: thumbs, ratings, "report a problem" with reasons.
- Implicit: accept, edit distance, regenerate, copy, abandon.
- Controls: tone, length, language, sources to use, memory on/off.
- Close the loop: show users that feedback changes results.

Feedback signals ranked by usefulness

Signal	Meaning	Use
Accept / edit / reject	Behavioural, honest	Best signal for quality; track edit distance
Thumbs up / down with reasons	Explicit	Triage failures; label evals
Regenerate or retry	Implicit dissatisfaction	Find weak prompts
Copy, share, send	Implicit success	Value metric
Escalation to human	Failure or trust gap	Measure and reduce
Free-text comments	Rich but sparse	Qualitative themes

Key takeaways

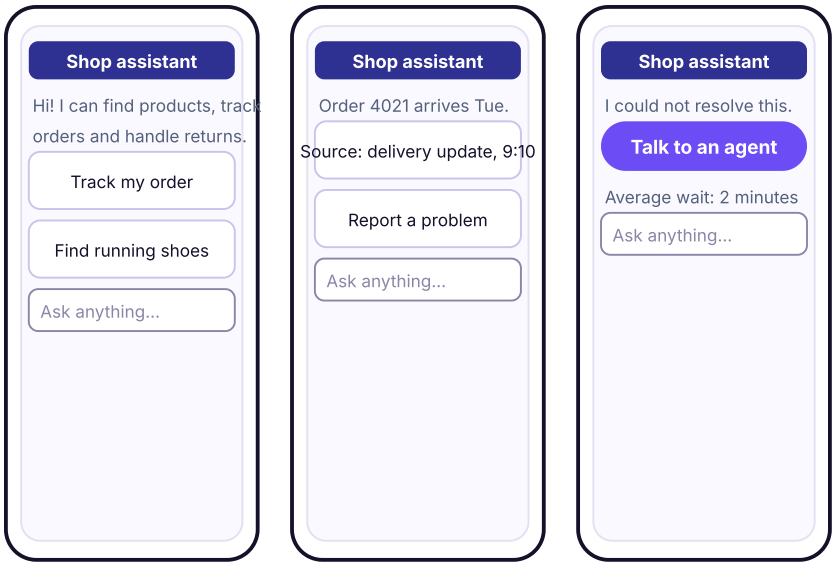
- Implicit signals are richer than thumbs.
- Give users meaningful controls.
- Feed signals into evals and improvements.

Practice exercise: List 6 feedback signals for an AI resume writer and how you'd use each.

INTERMEDIATE · CLASS 5

Conversational UI design in detail

Design chat and voice experiences people enjoy and trust.



Three states of a support assistant: start, answer with source, human handover

- **Opening:** greet briefly, state scope and examples, offer 3 to 4 tappable starters based on context (current order, current page).
- **Turn-taking:** stream answers so users see progress; keep paragraphs short; allow stop and edit of the last message.
- **Structure:** mix text with cards, buttons and forms for actions (track order, pay, reschedule). Chat is for discovery; structured UI is for decisions and payments.
- **Tone:** consistent, plain, respectful; mirror the user's language (English, Hindi, Hinglish); avoid fake human claims.
- **Memory and context:** show what is remembered, allow clearing; reuse order and profile data with permission.
- **Clarifying questions:** ask one at a time, offer choices.
- **Voice:** confirm critical details by reading them back; allow barge-in; display a transcript.
- **Exit paths:** always a way to reach a human, search, or the normal app screens.

Key takeaways

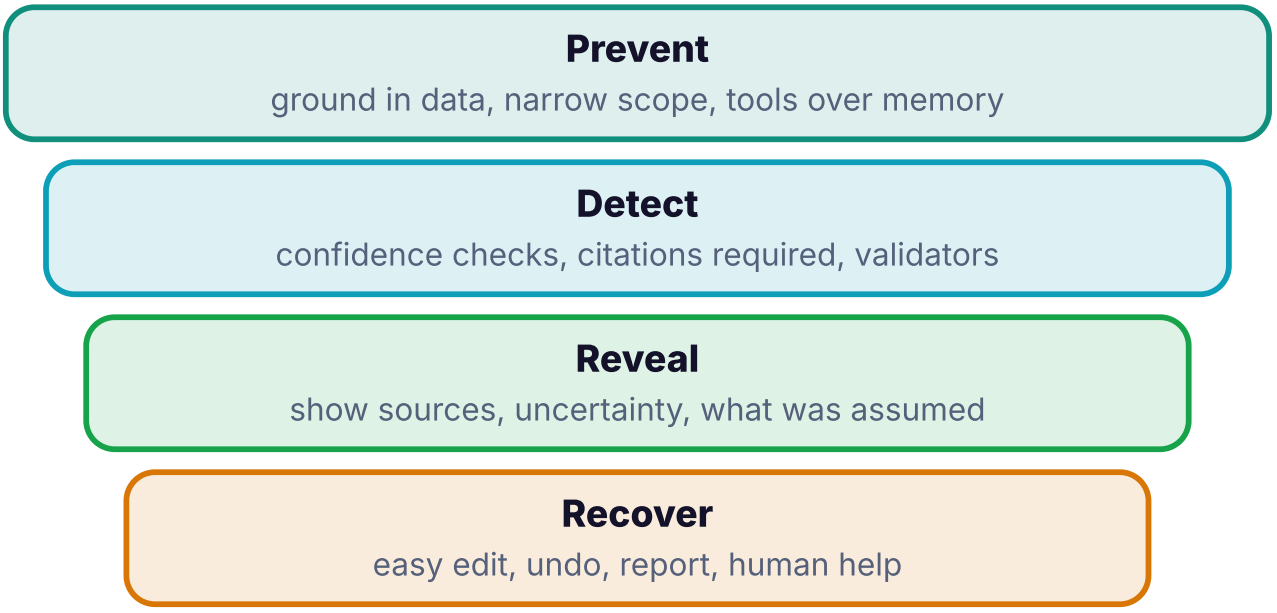
- Chat for exploration, structured UI for actions.
- State scope and give starters.
- Always provide a human or non-AI exit.

Practice exercise: Design the first screen and three starters for an AI assistant in an app you know.

INTERMEDIATE · CLASS 6

Designing for uncertainty, hallucination and verification

Make errors visible, recoverable and rare.



Four layers against AI mistakes

Technique	How it looks	When to use
Citations	Quoted passage with link	Knowledge answers
Abstention	'I could not find this in the policy. Contact support.'	Low-confidence or missing data
Show assumptions	'I assumed delivery to 560001'	Parameters guessed by AI
Verification step	'Check the amount before paying'	Money, health, legal
Dual display	AI text plus the underlying data table	Analytics and reports
Confidence language	'Likely', 'based on 3 sources'	Probabilistic answers
Review queue	Human checks before sending	High-stakes outputs

Avoid fake precision ("97% sure") unless the number is calibrated. Prefer specific evidence the user can check.

Key takeaways

- Layer prevent, detect, reveal and recover.
- Prefer abstaining to inventing.
- Add verification where stakes are high.

Practice exercise: List five likely errors of an AI feature you design and the UI response for each.

INTERMEDIATE · CLASS 7

AI UX in e-commerce and fintech: patterns and cases

See practical designs for shopping, money and support.

Use case	Good AI UX	Anti-pattern
Shopping assistant	Asks 2 questions, shows 3 products with reasons and price, filters link to normal search, can compare	Dumps 20 items in chat; no link to the catalogue; hides sponsored results
Review summary	Pros and cons with counts and links to reviews	One glowing paragraph with no evidence
Size and fit advice	Shows basis (your past purchases, brand size chart) and lets you override	'Perfect fit!' with no explanation
Spend insights (fintech)	Charts first, AI explanation second, with a drill-down to transactions	AI statements with no numbers
Payment help bot	Gets transaction status from the system, shows the actual status and next step	Guesses from general knowledge
Loan or investment guidance	Plain risks, no promises, disclaimers, human option	Personalised advice without regulated process
Support triage	Detects intent, shows relevant help, fast human handover with context	Loops of 'I did not understand'

Case: 'Where is my order?' assistant

1. Pull live order status from the order system (do not let the model guess).
2. Answer in one sentence with ETA and reason for delay.
3. Offer actions: notify me, change address (if allowed), cancel (if allowed), talk to an agent.
4. Measure: contact resolution without agent, CSAT, repeat contacts within 24 hours, wrong-status reports (target near zero).

Key takeaways

- Connect AI to systems of record for facts.
- Show reasons and evidence.
- Keep sponsored content and regulated advice clear.

Practice exercise: Redesign a support flow with AI and list the metrics and failure responses.

ADVANCED · CLASS 8

Failure, safety and handoff design

Design what happens when AI can't or shouldn't answer.

Situation	Design
Low confidence / no source	"I couldn't find this in your documents" + search or human
Sensitive request	Safe response + helpline/expert; no judgement
Agent about to act	Preview + confirm; clear undo window
Slow response	Streaming, progress steps, cancel button
Outage	Manual fallback path; status message

Failure types and designs

Failure	Example	Design response
Wrong answer	Invented policy	Cite or abstain; easy report; correct and learn
Refusal	Safety filter blocks a valid request	Explain why and offer an alternative
Misunderstanding	Intent not captured	Clarifying question; show interpretation
Slow or down	Timeout	Progress, partial results, fallback to search or human
Harmful output	Biased or unsafe text	Filters, quick takedown, escalation
Over-confident tone	Certain but wrong	Calibrated language; uncertainty cues

Key takeaways

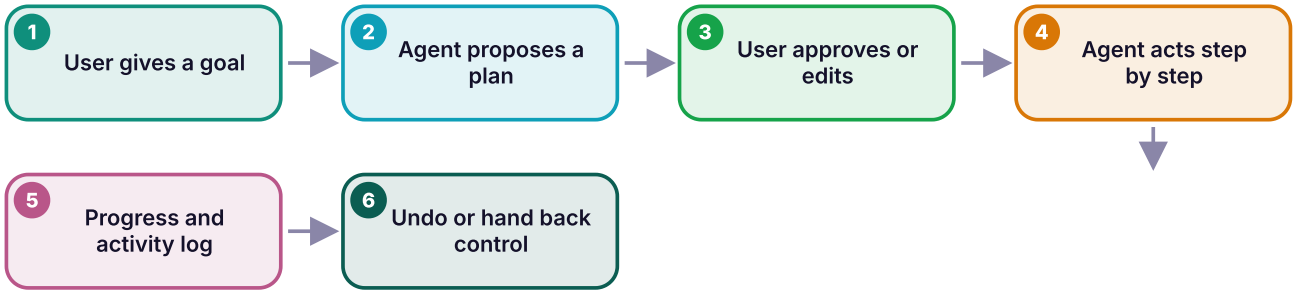
- Every AI feature needs a fallback path.
- Preview and confirm before actions.
- Stream and show progress for long tasks.

Practice exercise: Write failure states for an AI booking agent.

ADVANCED · CLASS 9

Agentic UX: autonomy, approvals and trust

Design AI that takes actions on the user's behalf.



A safe agent interaction loop

Autonomy level	Behaviour	Use for
Suggest	AI recommends; user acts	High stakes, early trust
Draft	AI prepares; user approves	Emails, orders, edits
Act with confirmation	AI performs after a yes per step or per batch	Bookings, refunds under a limit
Act and notify	AI acts within rules and tells the user	Low-risk repetitive tasks
Autonomous	Acts under policy with monitoring	Rarely; only with strong evals and limits

- **Permissions:** let users grant scopes (read calendar, spend up to Rs 500) and revoke easily.
- **Visibility:** activity log with who did what, when, and why; notifications for important actions.
- **Interruption and recovery:** pause, stop, undo, roll back; clear state when the agent fails halfway.
- **Plan previews:** show steps and expected effects before execution.
- **Limits:** spend caps, rate limits, allowed domains, human approval for irreversible actions.

Key takeaways

- Match autonomy to risk and trust.
- Preview, log, and make actions reversible.
- Give users scoped permissions and kill switches.

Practice exercise: Design the approval and undo flow for an AI agent that can rebook travel.

ADVANCED · CLASS 10

Research, evaluation and metrics for AI UX

Prove that an AI experience works for real users.

Method	What it gives you
Wizard of Oz	A human plays the AI to test ideas before building
Prompt-based prototypes	Fast real-model prototypes to see actual behaviour
Golden set with rubric	50 to 200 realistic tasks scored on accuracy, helpfulness, safety, tone
Failure taxonomy	Categorise errors to find the biggest fixes
Usability testing with think-aloud	See trust, confusion and recovery behaviour
Diary studies and logs	Real usage over time
A/B tests with guardrails	Impact on business metrics and safety

Metrics

Type	Examples
Outcome	Task success, time saved, resolution rate, conversion
Quality	Accept rate, edit distance, error reports, factuality score
Trust	Trust rating, verification clicks, repeat use, reliance (appropriate vs over-reliance)
Safety	Harmful output rate, escalations, policy violations
Efficiency	Latency, cost per task, human review rate

Key takeaways

- Test with real tasks and rubrics.
- Track acceptance and edits, not only likes.
- Measure appropriate reliance.

Practice exercise: Create a 20-task golden set and a scoring rubric for your AI feature.

ADVANCED · CLASS 11

Inclusive, multilingual and voice-first AI design

Make AI work for everyone, including Indian-language users.

- **Languages:** support English, Hindi and key regional languages, code-mixed Hinglish; test with real users; do not translate blindly.
- **Voice:** accents, noisy environments, speech disfluency; show transcripts; allow correction by typing.
- **Accessibility:** screen reader support for chat, keyboard access, captions, text size, no reliance on colour, clear focus order.
- **Literacy and numeracy:** short sentences, icons, audio replies, rupee formatting and amounts read out clearly.
- **Safety for children and vulnerable users:** age-appropriate content, parental controls where relevant, mental-health escalation paths.
- **Cultural fit:** names, addresses, festivals, units, date formats.
- **Bias testing:** check quality across languages, dialects, genders and regions.

Check	Pass criterion
Same task in Hindi and English	Success rates within an agreed gap
Voice in noise	Word error rate threshold; graceful retry
Screen reader run-through	All actions reachable and announced
Low-end device test	Response within target on a 3 GB phone and 4G

Key takeaways

- Test in every supported language with real users.
- Voice needs transcripts and correction.
- Measure quality gaps across groups.

Practice exercise: Write the test plan for a Hindi and English voice assistant.

ADVANCED · CLASS 12

Case studies: three AI experiences redesigned

See before/after reasoning for common AI features (illustrative).

Case 1: Support chatbot that frustrates users

Before	Problem	After
Greeting says 'Ask me anything'	Users ask out-of-scope questions	Scope message plus four starters
Long paragraphs	Hard to scan on mobile	Short answer plus buttons for actions
Guesses order status	Wrong ETA, angry users	Reads live status from the order system and cites time
No handoff	Loops of 'I did not understand'	After 2 misses show 'Talk to an agent' with chat context passed

Result targets: resolution without agent 35% to 55%; repeat contact within 24 hours down 20%; wrong-status reports near zero.

Case 2: AI writing assistant in a CRM

- Show a draft in an editable box with 'Shorter', 'Friendlier' and 'Add details' chips; never auto-send.
- Show which fields were used ('Based on the last email and deal stage').
- Track accept rate, edit distance, time per email; keep an 'off' switch.

Case 3: Shopping assistant

- Ask two questions (budget, use), show three products with reasons and a link to full search; label sponsored items; allow compare.
- Metrics: add-to-cart rate, search exits, return rate, trust score.

Key takeaways

- Fix scope, facts, structure and handoff first.
- Keep drafts editable and visible.
- Label sponsored content.

Practice exercise: Redesign one AI feature you use with the before/problem/after table.

ADVANCED · CLASS 13

Interview practice: AI UX scenario questions

Practise answers on trust, errors and metrics.

Q. Design an AI assistant for a banking app.
Model answer: I would pick a user job such as 'understand my spending and fix failed payments', not open-ended advice. The assistant reads live data from systems of record, shows numbers first and explanation second, offers actions through structured UI, adds verification for money actions, and hands off to humans for disputes. Metrics: task success, resolution without agent, trust rating, escalation and error reports.
Why it works: It scopes, grounds, structures and measures.

Q. Users do not trust our AI summaries. What do you do?
Model answer: I would add citations to the source, show what was included or excluded, let users expand the full content, track user edits and verification clicks, and test whether citations raise reliance. I would also improve accuracy through evals; UI cannot fix wrong summaries.
Why it works: It addresses both UX and quality.

Q. How would you measure whether an AI feature is helping?
Model answer: Outcome metrics (time saved, task success), quality (accept and edit rates, factuality score), trust (rating, appropriate reliance), safety (harm reports) and cost per task, compared with a control group.
Why it works: A balanced scorecard with a control.

Q. How much autonomy should an agent have?
Model answer: It depends on risk and reversibility. Start with suggestions, then drafts needing approval, then limited autonomous actions with caps, logs and undo, expanding as evals and user trust improve.
Why it works: A ladder tied to risk.

Key takeaways

- Scope the job, ground in data, show evidence.
- Balance outcome, quality, trust and safety metrics.
- Link autonomy to risk.

Practice exercise: Answer all four questions in writing, then compare with the model answers.

ADVANCED · CLASS 14

Testing AI UX and final project

Validate AI experiences with users and evals.

- Wizard-of-Oz tests before building (a human plays the AI).
- Prototype with real model outputs, including bad ones; watch how users react to errors.
- Measure: task success, time saved, trust calibration (do users catch wrong answers?), retention.
- Accessibility and Indian language support.

Final project

Design an AI feature end to end: pattern choice, key screens (Figma), expectation-setting copy, citation UI, feedback signals, failure states, and a 5-user test plan.

A simple AI UX test protocol

1. Choose 10 real tasks per user group.
2. Run them on a prototype with a real model or Wizard of Oz.
3. Record success, edits, time, trust rating and surprises.
4. Tag failures and fix by cause: prompt, data, UI, policy.
5. Re-run weekly on the same set to see progress.

Key takeaways

- Test with imperfect outputs.
- Measure trust calibration, not only satisfaction.
- Design and evals go together.

Practice exercise: Complete the AI UX final project.

TOOLKIT

Templates

AI UX checklist

Field	Fill in
Pattern chosen & why	
Expectation-setting copy	
Citation / explanation UI	
Controls (tone, length, sources)	
Feedback signals	
Failure & handoff states	
Accessibility & languages	
Test plan	

TOOLKIT

Cheat sheet

Term / tool	Meaning / use
Calibrated trust	Users trust AI as much as it deserves
Generate-then-edit	AI drafts, user polishes
Wizard of Oz	Human simulates AI in tests
Implicit feedback	Signals from behaviour (edits, accepts)
Layers	Prevent, detect, reveal, recover
Autonomy ladder	Suggest, draft, confirm, notify, autonomous
Best signal	Accept, edit or reject behaviour
Always provide	A human or non-AI exit

Glossary

Handoff	Passing a task to a human
Streaming	Showing output as it's generated
Hallucination	Confident but false AI output
Grounding	Tying answers to trusted data
Abstention	Choosing not to answer when unsure
Wizard of Oz	A human simulating AI in a test
Golden set	Curated tasks used to score quality
Autonomy level	How much the AI may do without approval
Appropriate reliance	Trusting AI when it is reliable and not when it is not
Barge-in	Interrupting a voice assistant while it speaks
Code-mixing	Mixing languages in one sentence, like Hinglish
Edit distance	How much a user changes an AI draft
Escalation	Handing over to a human

INTERVIEW PREP

15 interview questions with answers

Q1. Key principles of AI UX?
Set expectations, design for uncertainty, transparency with sources, user control and recovery, graceful failure with handoff.

Q2. When not to use chat?
When a task is small and in a workflow; an inline button or suggestion is faster and clearer.

Q3. How do you build trust?
Citations, plain explanations, labels for AI content, consistent quality, and easy correction.

Q4. What feedback signals matter?
Accept/edit/regenerate/abandon rates and edit distance, plus explicit ratings with reasons.

Q5. How do you test AI UX?
Wizard-of-Oz, prototypes with real and bad outputs, task success and trust calibration metrics.

Q6. Design for an AI agent action?
Preview, confirmation, progress, undo window, and clear logs.

Q7. What is different about designing AI interfaces compared with traditional software?
Outputs are probabilistic and open-ended, so design must set expectations, show sources, handle errors, give control and gather feedback, rather than assume the same result every time.

Q8. How do you design for hallucinations?
Prevent with grounding and narrow scope, detect with checks and citation requirements, reveal with sources and uncertainty, and recover with easy edits, reporting and human help.

Q9. When is chat the wrong interface?
When tasks are structured and repeatable (payments, forms, comparisons). Use structured UI, with chat as an optional helper.

Q10. How do you build trust in AI features?
Be transparent about limits and data use, cite sources, keep users in control, show consistent quality, and make failure recovery easy.

Q11. What metrics would you use for an AI assistant?
Task success, resolution without human, accept and edit rates, CSAT, trust, escalation rate, factuality and safety scores, latency and cost per task.

Q12. What is Wizard of Oz testing?
A human secretly acts as the AI behind a prototype so you can test the experience before building the model.

Q13. How do you decide how autonomous an agent should be?
Based on risk, reversibility, strength of evals and user trust; start with suggestions and drafts, then add approvals and limited autonomous actions.

Q14. How do you handle AI refusals gracefully?
Explain briefly why, avoid blame, offer a safe alternative or human help, and monitor over-refusal.

Q15. How would you support Hinglish users?
Test with real code-mixed inputs, support mixed-script text, accept voice, evaluate quality by language, and avoid forced translation.

LEARN MORE

Further reading and tools

- Google PAIR Guidebook: https://pair.withgoogle.com/guidebook/
- Microsoft HAX Toolkit (Human-AI guidelines): https://www.microsoft.com/en-us/haxtoolkit/
- Nielsen Norman Group: AI articles: https://www.nngroup.com/topic/ai/
- Shape of AI (UX patterns): https://www.shapeof.ai/
- Google PAIR: People + AI Guidebook: https://pair.withgoogle.com/guidebook/
- Microsoft HAX Toolkit: https://www.microsoft.com/en-us/haxtoolkit/
- Apple HIG: machine learning: https://developer.apple.com/design/human-interface-guidelines/machine-learning

© SkillNavora · skillnavora.in · Part of the AI Product Management and GenAI Launchpad programs. Practise with the free AI Tutor, AI PM Toolkit and Practice Lab on skillnavora.in.